


# APACHE FLINK CLOUD BURSTING

Anwasha Saha, Sakshi Sharma, Sanath Bhimsen,  
Showndarya Madhavan, Ye Tian, Yujie Yan



# Project goal

The goal of this project is to design and implement an adaptive Flink application that leverages the “cloud bursting” technique as an alternative to back-pressure.



---

## Main achievements

- Implemented an Fault tolerant adaptive Flink application leveraging "cloud bursting" technique.
- Developed a controller that periodically collects execution metrics in order to invoke cloud bursting when workload variations are detected.
- Integrated with AWS lambda functions written in Python and Java with code package deployment.
- Implemented a custom source with dynamic input rate and spikes generation.
- Tested and evaluated the system's performance under different workload conditions to demonstrate its effectiveness in handling workload spikes and as an alternative to back-pressure.
- Documented the implementation details and provided clear instructions for reproducing the results.

---

# Demo

+



# Experimental Setup

Ran our experiments with **3** different pipelines

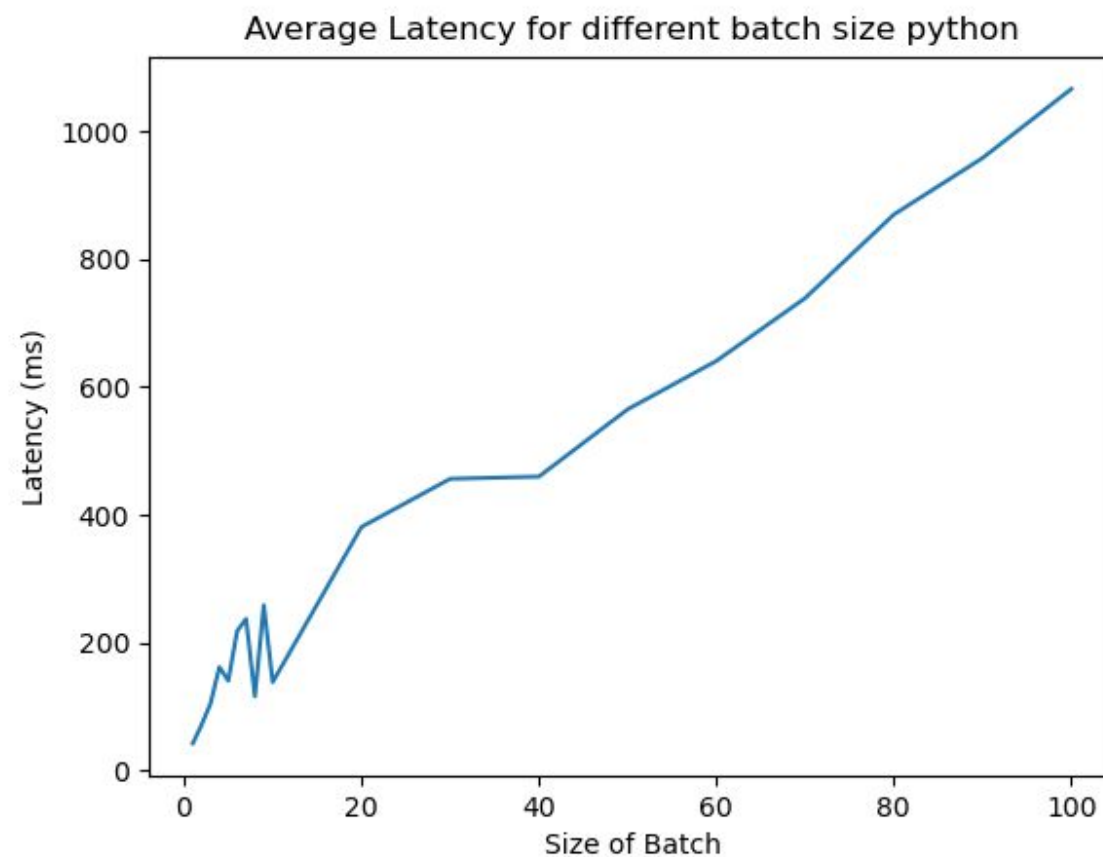

## Vanilla Pipeline



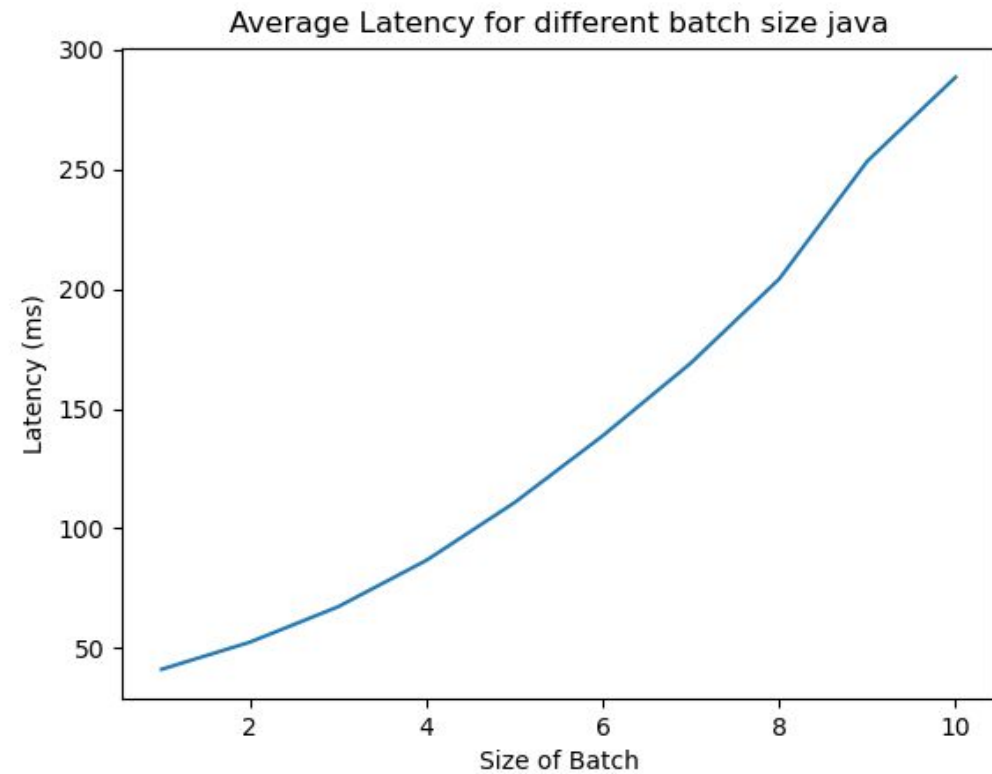
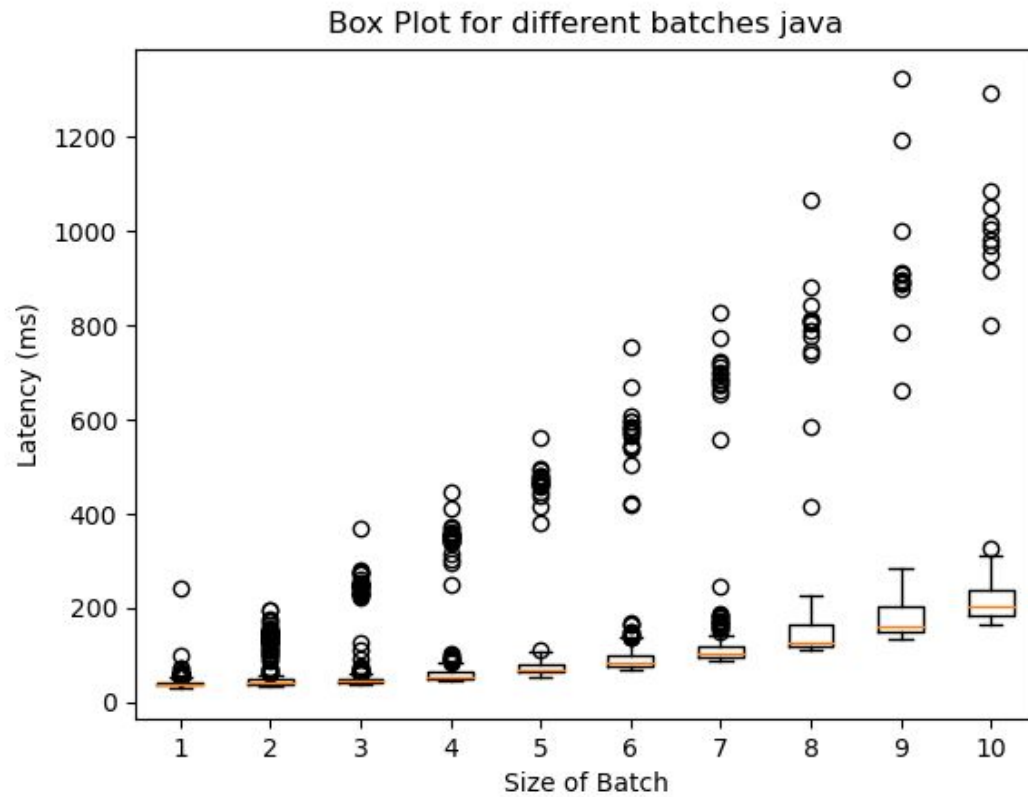
## Modified Pipeline without Lambda



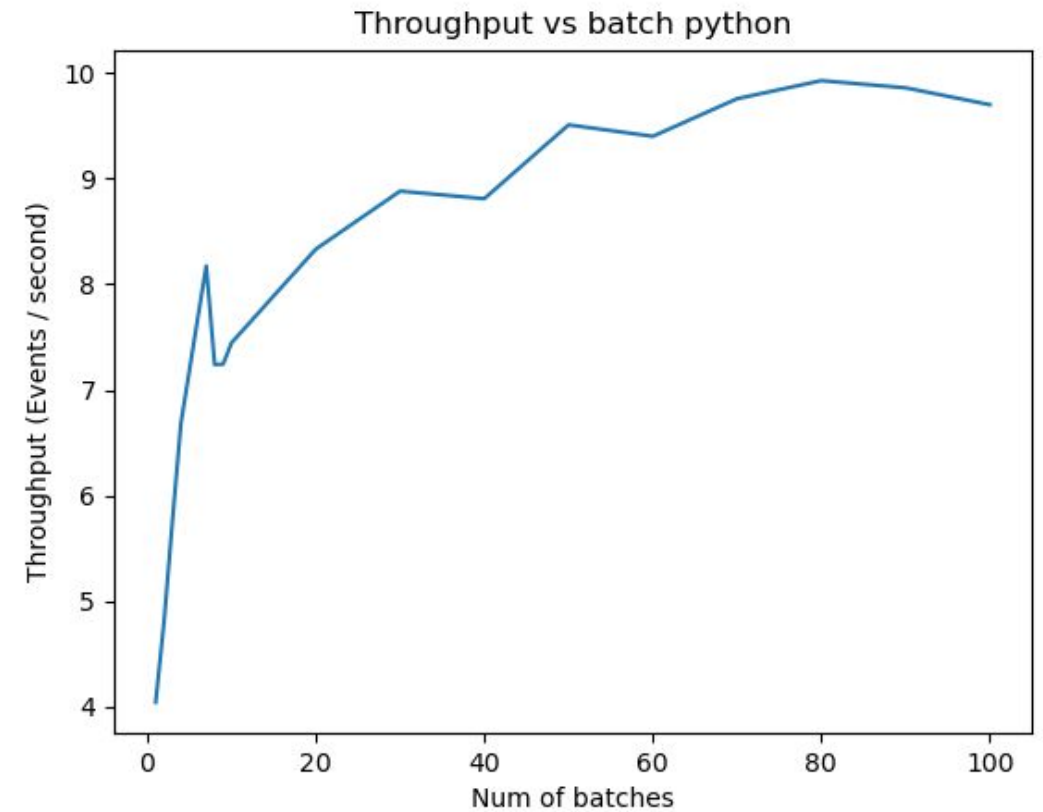
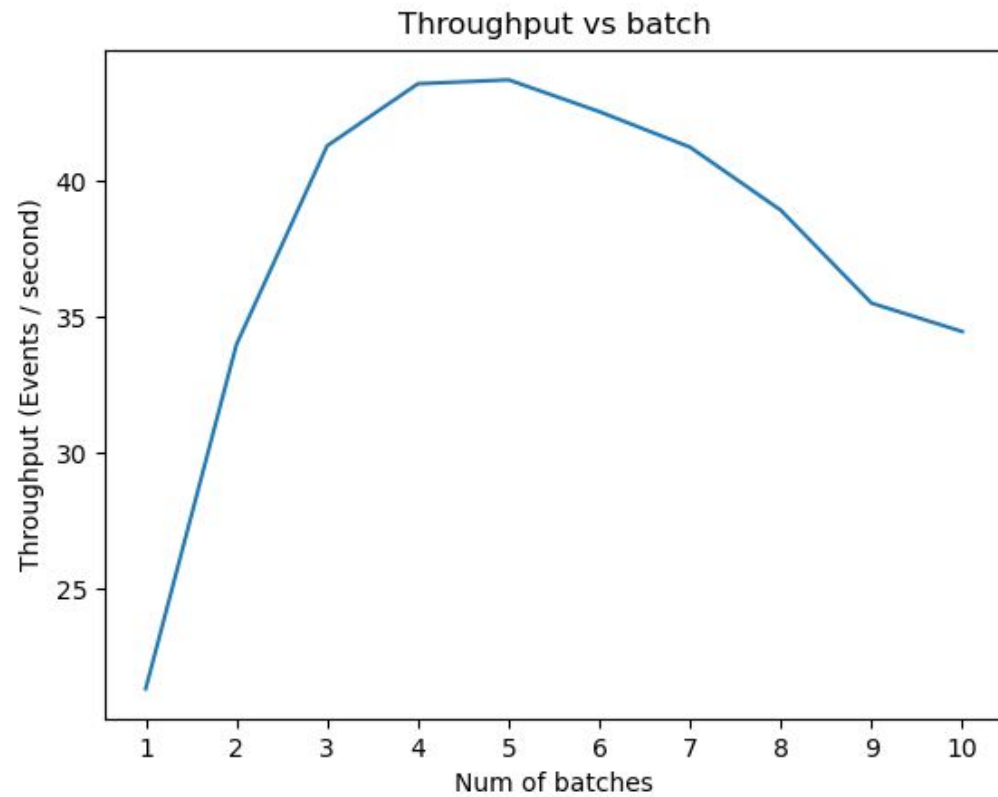
# Experimental results



# Experimental results



# Experimental results





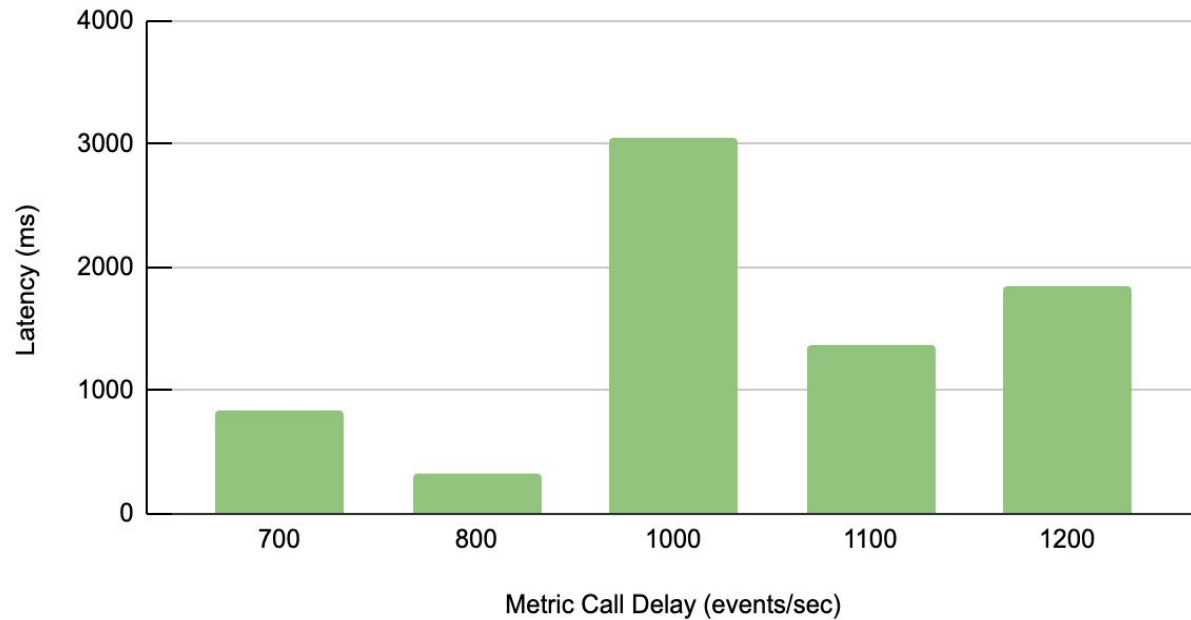
# Experimental results

Input Rate = 1000 events/s  
Duration of Experiment = 5mins  
Parallelism = 1  
Max Records Sent = unlimited,  
Batch size change-every 8 secs

Input Rate = 1000 events/s  
Duration of Experiment = 5mins  
Metric Call Latency = 800 ms  
Max Records Sent = 100000

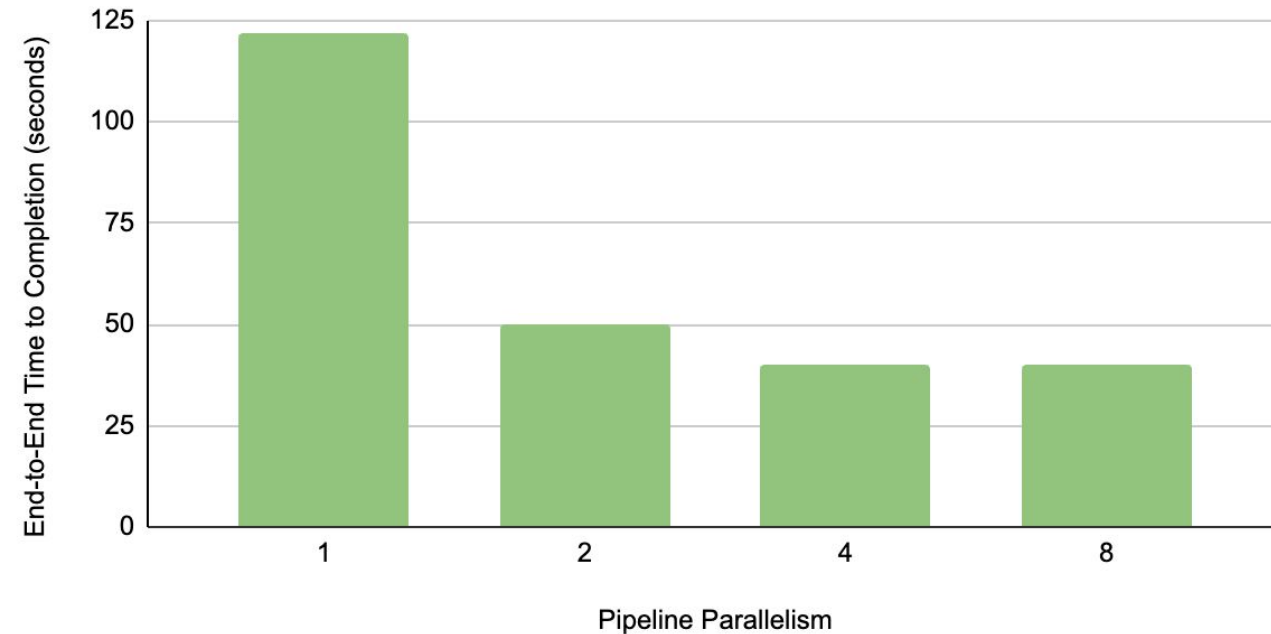
## Metric Call Delay vs Latency

Lambda Pipeline



## Parallelism vs Time to Completion

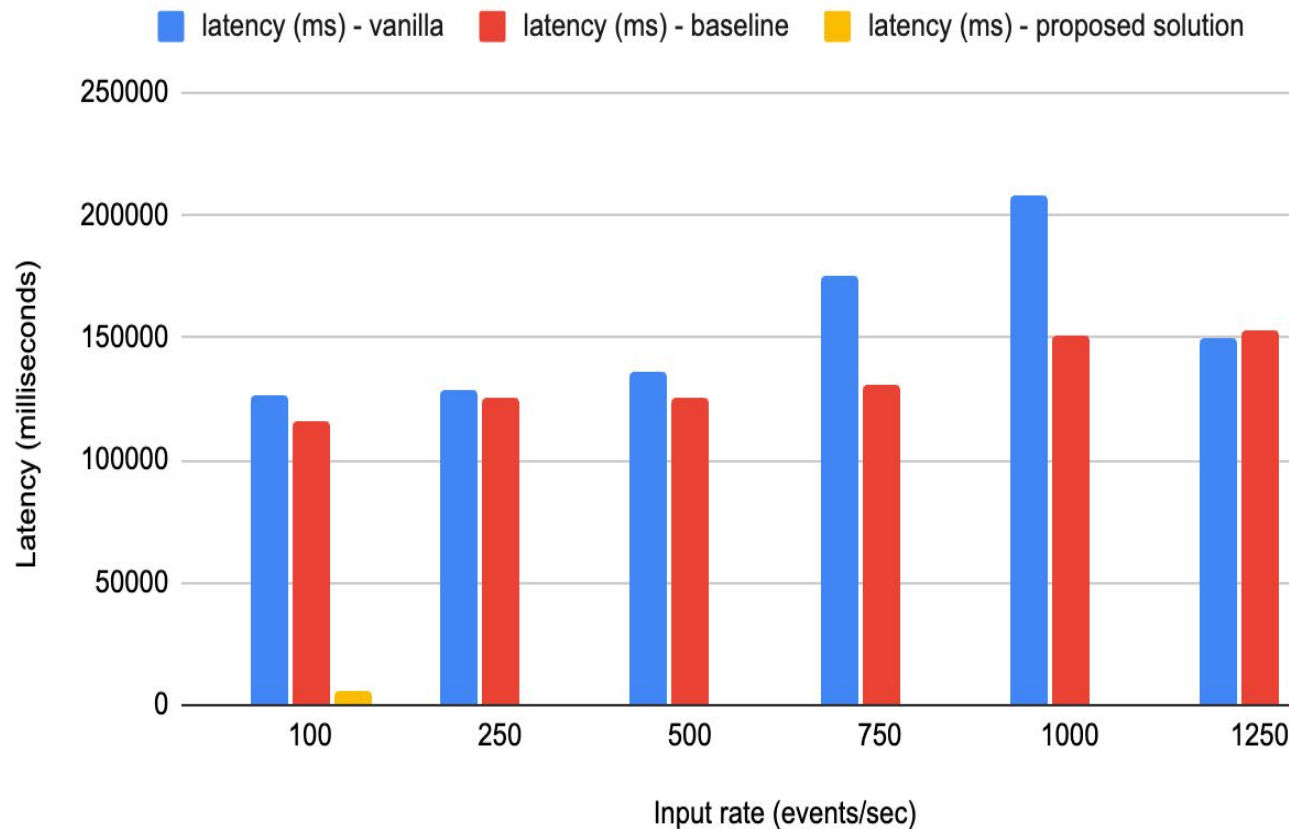
Lambda Pipeline



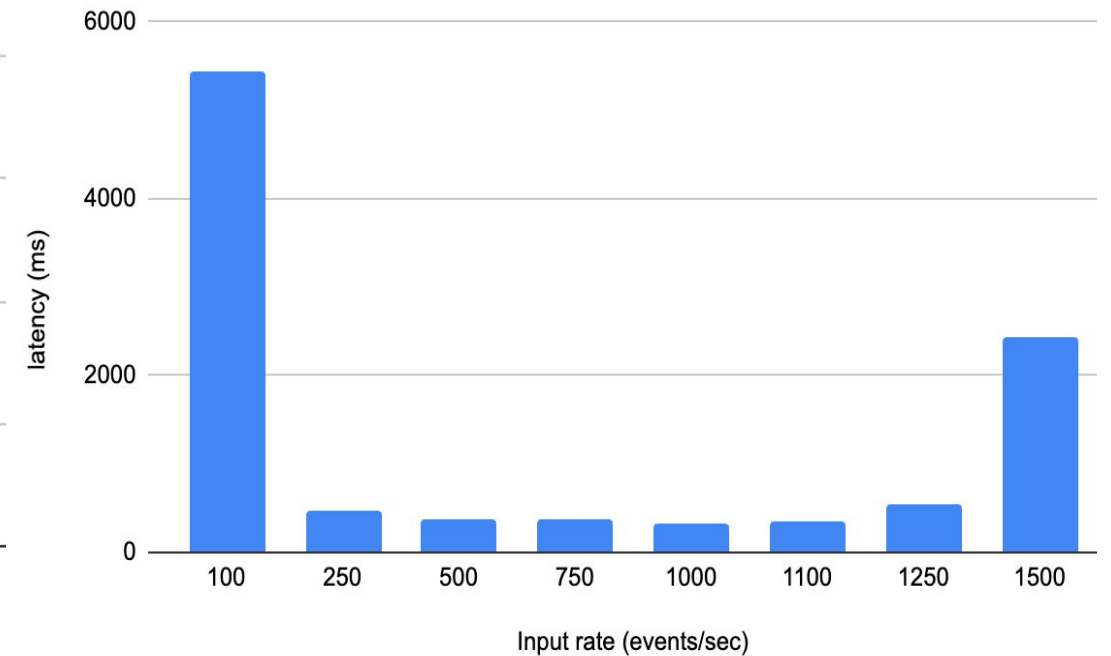
# Experimental results

## Comparison of all 3 pipelines

Vanilla vs Modified Pipeline (without lambda) vs Modified Pipeline (with lambda)



latency (ms) vs. Input rate (events/sec)



# Experience, Challenges and Future work

---

## Challenges

- Hyper parameter tuning
- Integration with lambda
- Policy metrics calculation
- Integration
- Flink documentation redirects to erroneous URL at many instances.

## Future work

- Handling stateful operations
- Security with lambda
- Fault tolerance
- Fetch the flink metrics directly from code instead of getting it from the flink web UI



## Key takeaways

Overhead increases a little, but we save on processing time and reduce back pressure on the pipeline



## Accumulated Backpressure time vs Event Rate

Duration of Experiment = 5mins  
Parallelism = 1  
Max Records Sent = unlimited,  
Batch size change-every 8 secs  
Metric Call Latency = 1000 ms

